

程序代写代做 CS编程辅导



proposal

April 9, 2023

1 CS 449 Final Proposal

Due: April 18, 2023 at 11:59pm

1.1 1. Names and Net IDs

List your group members

1.2 2. Abstract

Your abstract should be two or three sentences describing the motivation for your project and your proposed methods.

For example:

Our final project seeks to use a variety of `sklearn` models to classify handwritten digits in the MNIST dataset. We will compare models such as Logistic Regression and `DecisionTreeClassifier`.

1.3 3. Introduction

Why is this project interesting to you? Describe the motivation for pursuing this project. Give a specific description of your data and what machine learning task you will focus on.

For example:

It is very important for us to be able to automatically recognize handwritten digits so that the Postal Service can identify whether a letter has been sent to the correct address. We will use a large dataset of handwritten digits and train our models to input the black-and-white pixels of those digits and output the number that was written. [etc. etc.]

1.4 4a. Describe your dataset(s)

List the datasets you plan to use, where you found them, and what they contain. Be detailed! For each dataset, what does the data look like? What is the data representation? (e.g., what resolution of images? what length of sequences?) How is the data annotated or labeled? Include citations for the datasets. Include at least one citation of previous work that has used your data, or explain why no one has used your data before.

For example:

We use the popular MNIST dataset, which contains 70,000 different 28x28 pixel greyscale images of handwritten digits. This dataset was created by the American Census Bureau in 1998 ...


1.5 4b. Load your data

Demonstrate that you have made some progress with getting your dataset ready to use. Load at least a few examples as best you can

```
[1]: # For example:
import matplotlib.pyplot as plt

from sklearn.datasets import load_digits

X, y = load_digits(return_X_y=True)
fig, axes = plt.subplots(nrows=1, ncols=5, figsize=(20, 5), sharex=True,)
plt.gray()
for i, axis in enumerate(axes):
    axis.matshow(X[i, :].reshape(8, 8))
plt.show()
```



1.6 5. Methods

Describe what methods you plan to use. Cite at least one or two relevant papers. What kind of neural network architectures will you use? What loss function(s) will you use and why? How will you evaluate your model's ability to generalize to new data?

For example:

This is a standard supervised learning task, and we will use `sklearn`'s Logistic Regression model to predict digit labels from their pixels. The model will contain one weight per pixel. We will train our model using Cross-Entropy loss, because...

1.7 6. Deliverables

Include at least six goals that you would like to focus on over the course of the quarter. These should be nontrivial, but you should have at least one and hopefully both of your "Essential" goals done by the project update, due . Your "Stretch" goals should be ambitious enough such that completing one is doable, but completing both this quarter is unlikely.

1.7.1 6.1 Essential Goals

- (At least two goals here)

For example:

We will use a **LeNet** Multilayer Perceptron to train and test on our MNIST data.

1.7.2 6.2 Desired

- (At least two goals here)

For example:

We will conduct a thorough hyperparameter search for all our models and provide a detailed and practical guide for how a researcher should search for good hyperparameter values.

1.7.3 6.3 Stretch Goals

- (At least two goals here) > For example: >> We will conduct a manual analysis of the digits that our model gets wrong and use a GAN to create new images that help us learn a more robust classifier.

1.8 7. Hopes and Concerns

What are you most excited about with this project? What parts, if any, are you nervous about? For example:

For example:

We're worried that we'll get really bored of staring at pixelated hand-written digits for hours on end.

1.9 8. References

Cite the papers or sources that you used to discover your datasets and/or models, if you didn't include the citation above.

For example:

LeCun, Yann, et al. "Gradient-based learning applied to document recognition." Proceedings of the IEEE 86.11 (1998): 2278-2324.